

Artificial Neural Networks and Deep Learning

Week 5


Recurrent neural networks

Recurrent Neural Networks

THE neural network architecture to use for sequential data


Recurrent neural networks

> Ways to process sequential data



The **problem** with **all** feed forward neural networks

- Input and output must be of hard-assigned dimensions
- The network makes a fixed number of computations
- If input is sequence-like (video, sound, etc.) the network is ignorant to the order of samples




<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

The **problem** with **all** feed forward neural networks

> **Solution:** *Recurrence!*


- Input and output must be of hard-assigned dimensions
- The network makes a fixed number of computations
- If input is sequence-like (video, sound, etc.) the network is ignorant to the order of samples



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Recurrent neural networks

> Fundamental idea




$$h_t = f_W(h_{t-1}, x_t)$$

Notice: W is the same in each iteration

Recurrent neural networks

> Fundamental idea



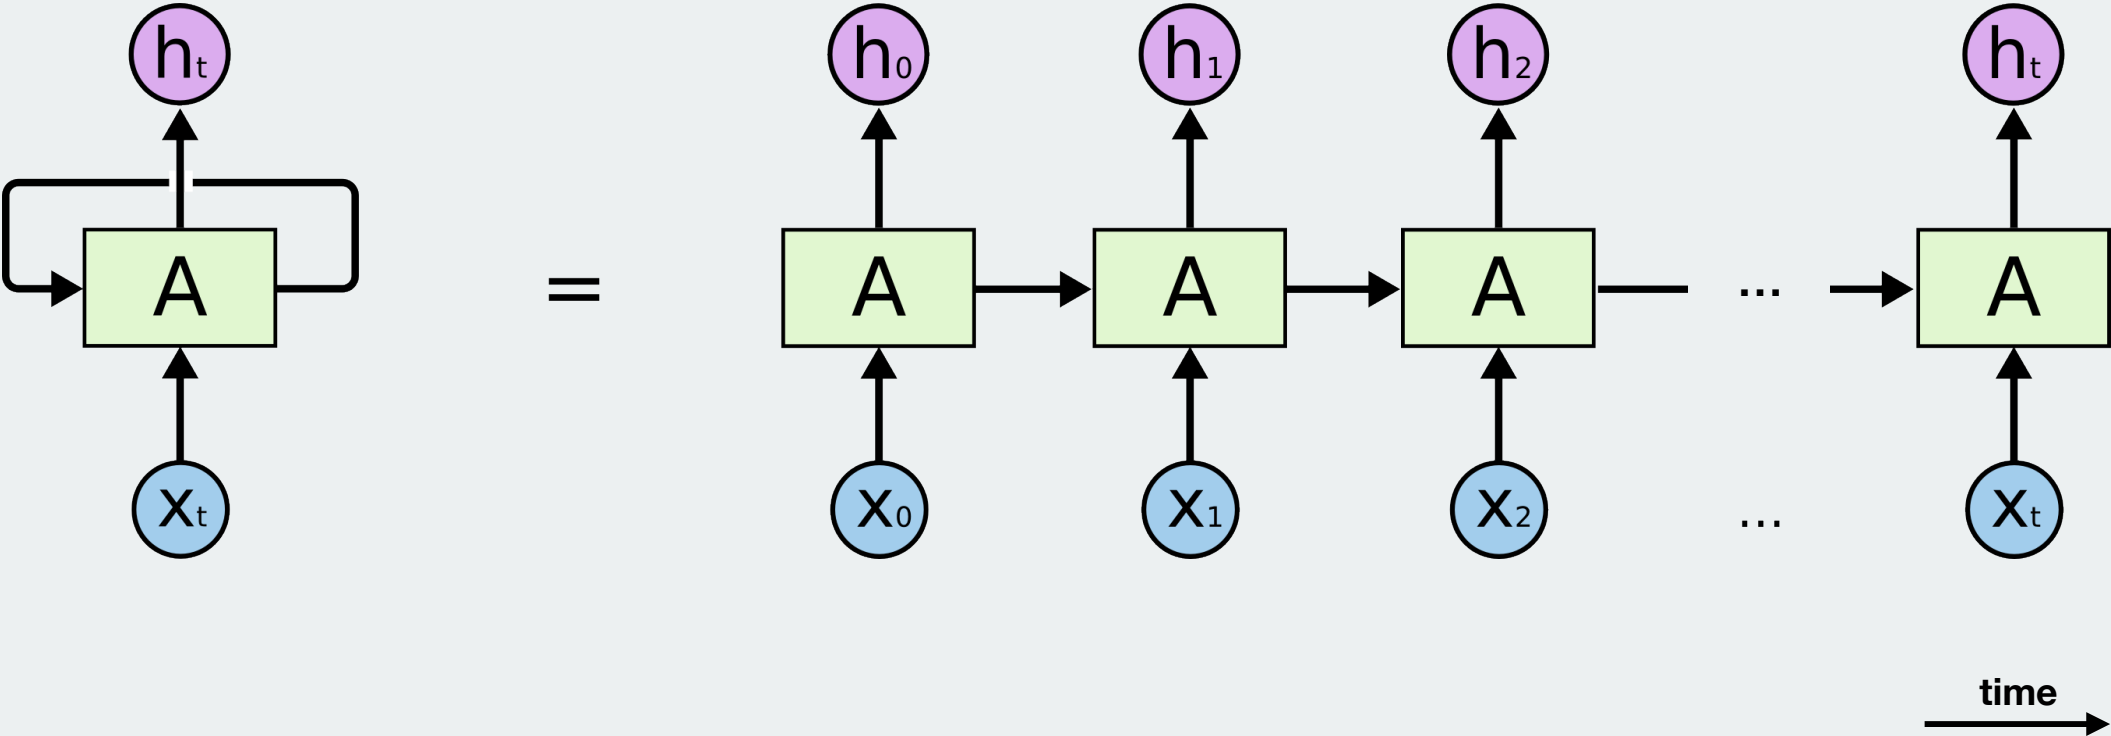
$$h_t = f_W(h_{t-1}, x_t)$$

$$y_t = W_{hy}h_t$$

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Recurrent neural networks

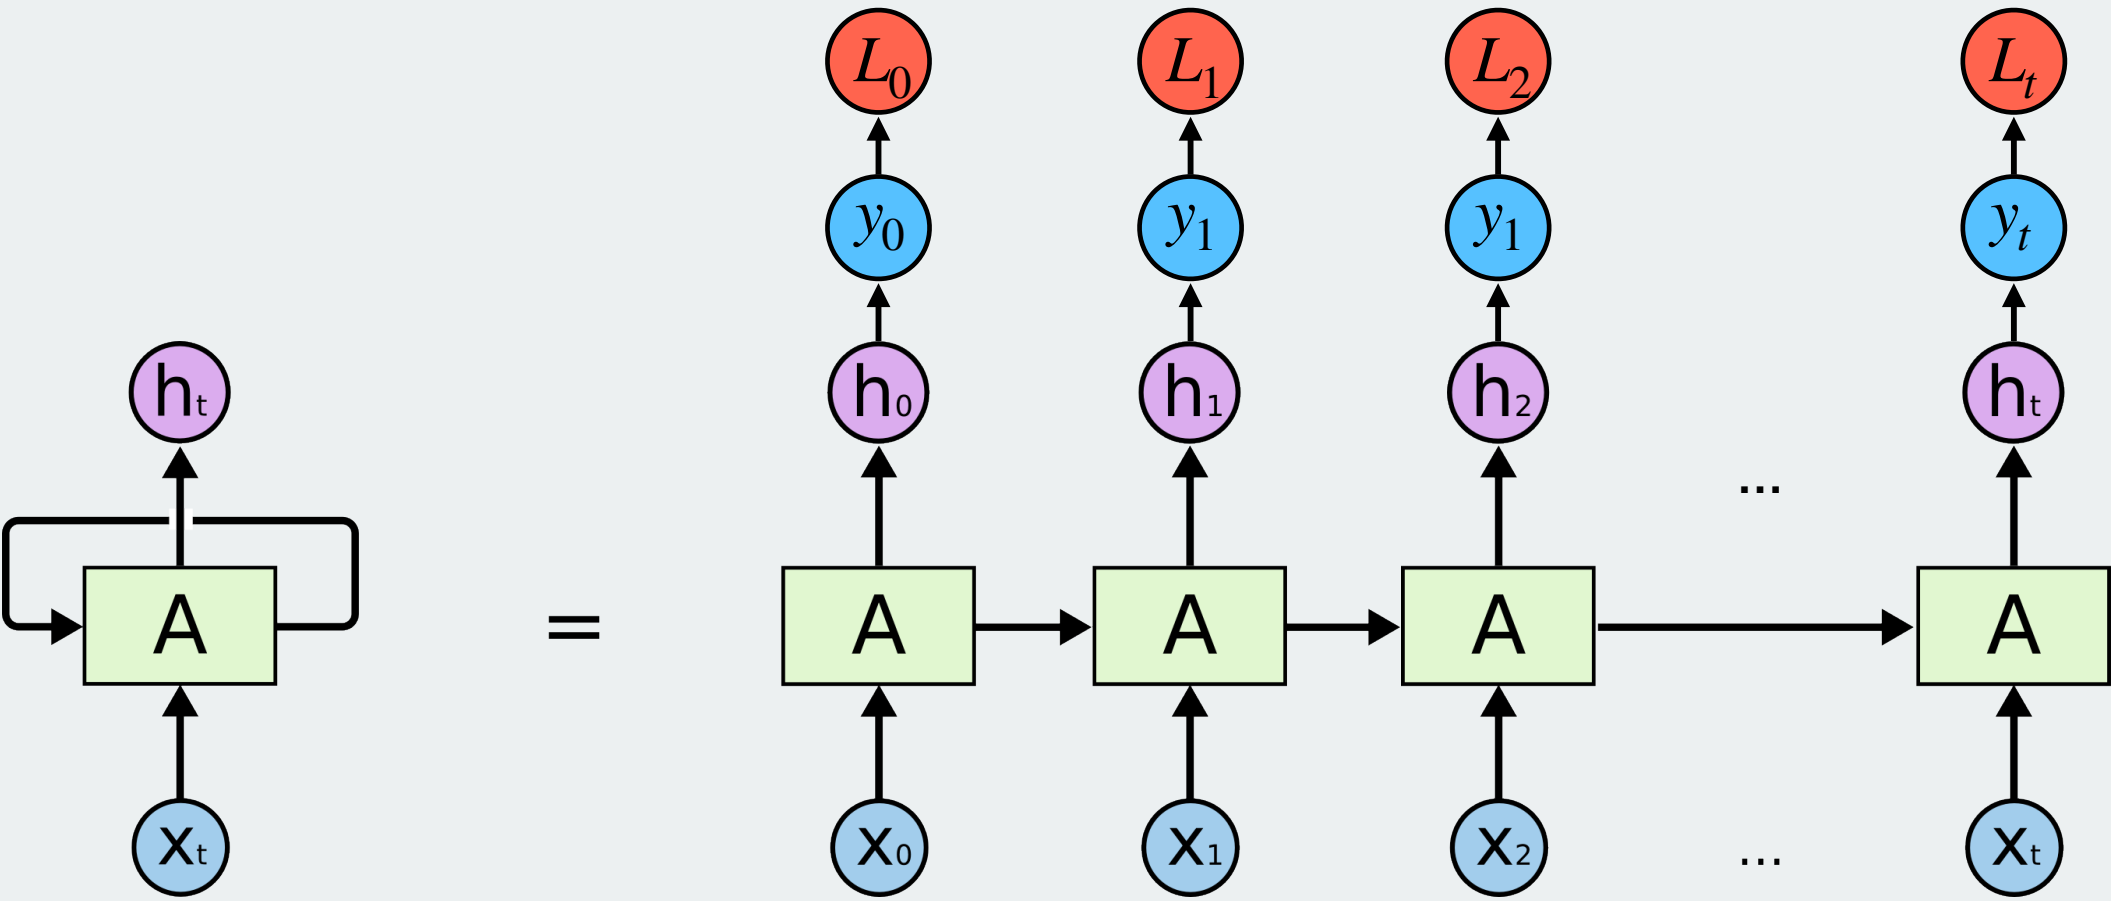
> Unrolled in time



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Recurrent neural networks

> Backpropagation

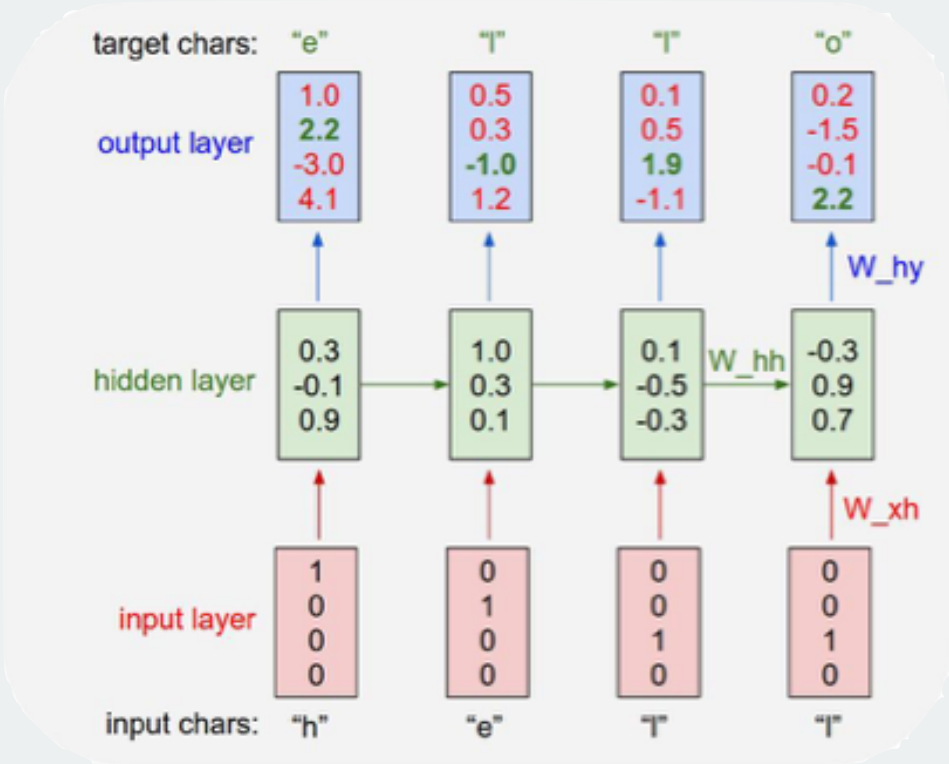


<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Recurrent neural networks


> Example: predicting next character

training sequence: “hello”



Recurrent neural networks


> Vanilla RNN, architecture



Recurrent neural networks


> Vanilla RNN, architecture

4: because dotting (h, t) onto it should result in a new vector with 4 elements
7: because the (h, t) vector we are dotting onto it has 7 elements in it




Recurrent neural networks

> Vanilla RNN, backpropagation



- > Vanilla RNN, backpropagation



Problem: Repeated multiplications by **W** during backpropagation

Leads to: Exploding/vanishing gradients

Solution: Gradient clipping (solves exploding gradients), or **change architecture**

Recurrent neural networks

> Vanilla RNN vs. LSTM

Vanilla RNN

$$h_t = \tanh \left(W \begin{bmatrix} h_{t-1} \\ x_t \end{bmatrix} \right)$$

Long Short Term Memory (LSTM)

$$\begin{bmatrix} i \\ f \\ o \\ g \end{bmatrix} = \begin{bmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{bmatrix} W \begin{bmatrix} h_{t-1} \\ x_t \end{bmatrix}$$

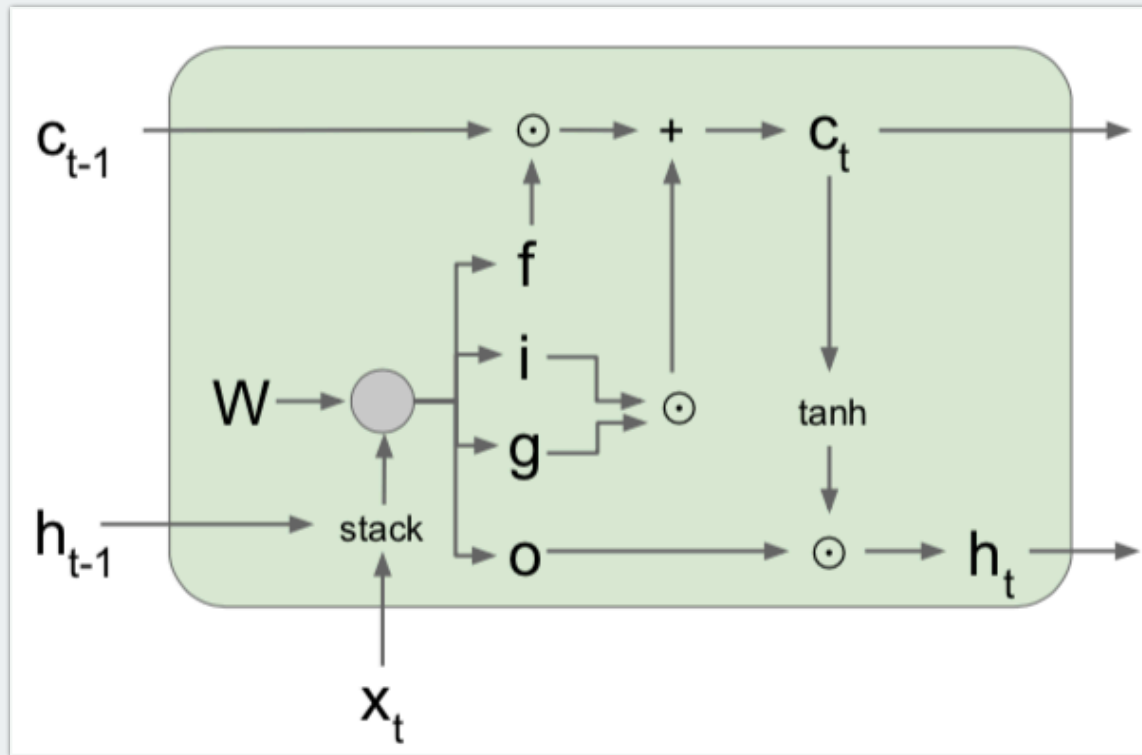
$$c_t = f \odot c_{t-1} + i \odot g$$

$$h_t = o \odot \tanh(c_t)$$

<http://cs231n.stanford.edu/syllabus.html>

Recurrent neural networks

> Vanilla RNN vs. LSTM



Long Short Term Memory (LSTM)

$$\begin{bmatrix} i \\ f \\ o \\ g \end{bmatrix} = \begin{bmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{bmatrix} W \begin{bmatrix} h_{t-1} \\ x_t \end{bmatrix}$$

$$c_t = f \odot c_{t-1} + i \odot g$$

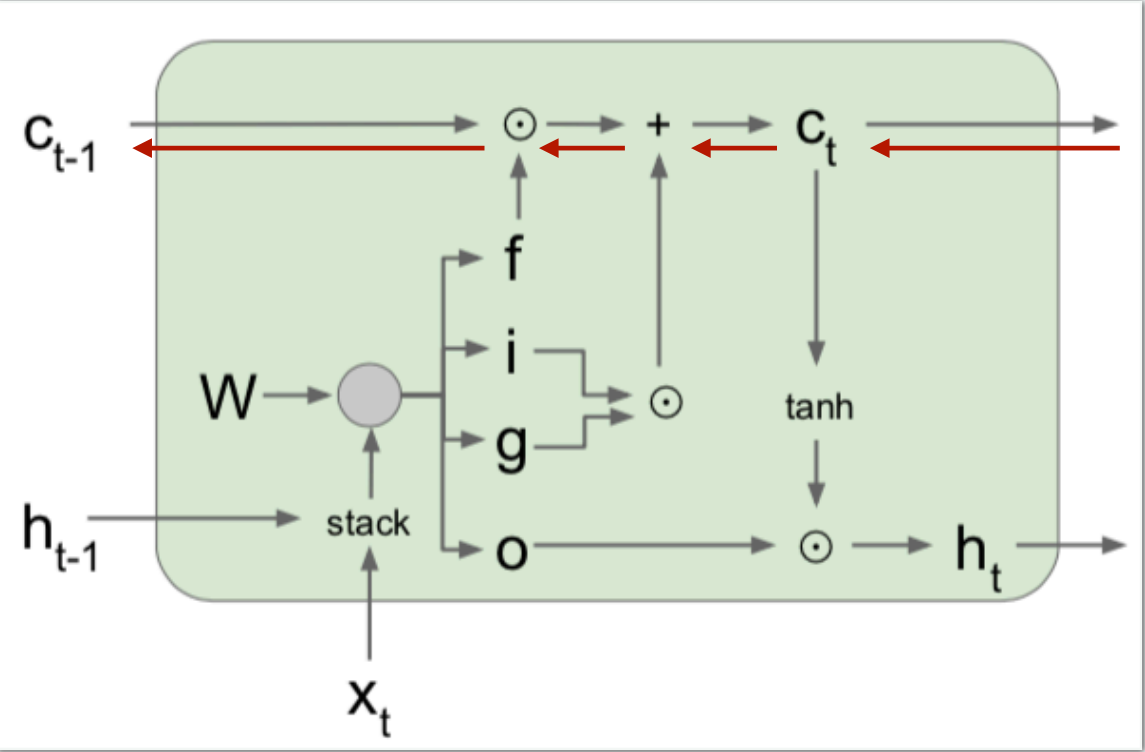
$$h_t = o \odot \tanh(c_t)$$

<http://cs231n.stanford.edu/syllabus.html>

Recurrent neural networks

> Vanilla RNN vs. LSTM

“Gradient highway”: solves vanishing/exploding gradient problems




Long Short Term Memory (LSTM)

$$\begin{bmatrix} i \\ f \\ o \\ g \end{bmatrix} = \begin{bmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{bmatrix} W \begin{bmatrix} h_{t-1} \\ x_t \end{bmatrix}$$
$$c_t = f \odot c_{t-1} + i \odot g$$
$$h_t = o \odot \tanh(c_t)$$

Recurrent neural networks

> Vanilla RNN vs. LSTM



Long Short Term Memory (LSTM)

$$\begin{bmatrix} i \\ f \\ o \\ g \end{bmatrix} = \begin{bmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{bmatrix} W \begin{bmatrix} h_{t-1} \\ x_t \end{bmatrix}$$
$$c_t = f \odot c_{t-1} + i \odot g$$
$$h_t = o \odot \tanh(c_t)$$